



Christopher R


ESSENTIAL FOR BUILDING AI AGENTS





Christopher R

WHY LEARN TO BUILD AI AGENTS?



AI agents go beyond simple chatbots. They combine reasoning, memory, and tools to act autonomously, making them capable of performing real tasks like research, scheduling, coding, or customer support.



They integrate LLMs with workflows and APIs.



They can learn and adapt over time.



They serve as the foundation for next-gen automation and intelligent systems

A shopping bot that finds the cheapest price
This guide gives you 12 step-by-step lessons
to design and scale powerful agents.






Christopher R

UNDERSTAND AGENT ARCHITECTURE

Before building, know the blueprint. Every AI agent has four pillars:



Frameworks to try:



Langchain



Haystack



AutoGen

Tip

Start with one task (e.g., answering FAQs with memory) before attempting multi-tool autonomous agents.





Christopher R

CHOOSE THE RIGHT LLM

Not all LLMs are equal. The right choice depends on task, cost, and infrastructure.



GPT-4o:

Great for reasoning + function calling.



Claude 3.5:

Long-context handling, great for summarization.



Mistral / LLaMA:

Great for reasoning + function calling.



Mixtral / Gemini:

Multi-modal and cutting-edge.

Tip

Run task-specific benchmarks (accuracy, latency, cost) before committing to one model.







Christopher R

ADD MEMORY FOR CONTEXTUAL INTELLIGENCE

Without memory, agents forget everything after each task.
You need both:



Long-term memory:
Keeps track of current
conversations (chat history).



Short-term memory:
Stores facts, documents, and
interactions.

Tools:



Pinecone



Weaviate



Milvus



Chroma

Tip

Separate factual memory (knowledge base) from episodic memory (conversations) for better control.





Christopher R

EQUIP AGENTS WITH TOOLS

Agents need actions, not just words. Tools extend their capabilities.


Types of Tools:



APIs
(weather, finance, CRM).



Databases
(SQL, NoSQL, vector stores).



Utility functions
(calculators, code execution).

Frameworks:



LangChain
Toolkits



OpenAI
Functions



n8n

Tip

Start with essential tools only—too many tools can confuse the reasoning engine.





Christopher R

MASTER PROMPT ENGINEERING

Prompts are the control system for your agent. Poor prompts = unpredictable outputs.

Techniques:



Chain of Thought: Make reasoning explicit.



ReAct Pattern: Alternate reasoning + action.



Role Assignment: "You are an expert planner..."

Frameworks:



PromptPerfect



FlowGPT libraries

Tip

Build a prompt library and reuse tested prompts across agents for reliability.





Christopher R

HANDLE MULTI-STEP REASONING

Most tasks require step-by-step execution, not single-shot answers.

Patterns to use:



ReAct: Think → Act → Observe → Repeat.



Plan-and-Execute: Create a roadmap, then act.



Tree of Thoughts: Explore multiple reasoning paths before choosing.

Tip

Always ask agents to output in structured formats (JSON, YAML) to reduce hallucinations.





Christopher R

INTEGRATE WITH WORKFLOWS

An agent is rarely enough—you need workflows to tie everything together.

Automation Platforms:



n8n: Open-source automation.



Make.com: Visual workflow builder.



Zapier AI: Business-focused integration.

Best Practice:

Insert human-in-the-loop steps for approvals in critical tasks.

Tip

Think of agents as specialized team members, not entire teams.





Christopher R

ENSURE OBSERVABILITY & MONITORING

You can't improve what you can't measure.



What to track: Inputs, outputs, tool usage, errors, reasoning traces.



Tools: LangSmith, Helicone, Phoenix, HoneyHive.



Why it matters: Helps debug bad outputs, monitor costs, and improve accuracy.

Tip

Build a feedback loop where users can rate or flag agent responses.





Christopher R

OPTIMIZE FOR COST & LATENCY

LLM calls are expensive. Optimize early.

Strategies:



Cache frequent queries with Redis/GPTCache.



Use smaller models for simple tasks.



Batch requests instead of calling APIs multiple times.

Tools:



LiteLLM for routing requests across multiple LLMs.

Tip

Track cost per task, not just per token—this shows the real efficiency.






Christopher R

ADD GUARDRAILS FOR SAFETY

Agents need boundaries to prevent mistakes or unsafe actions.



Guardrails: Restrict harmful topics, filter outputs, enforce formats.



HoneyHive

Tools: Guardrails AI, NeMo Guardrails, HoneyHive.



Methods: Policy enforcement, content moderation, sandboxing risky actions.

Tip

Always define “allowed actions” vs “forbidden actions” before deployment.





Christopher R

TEST & VALIDATE LIKE SOFTWARE

Agents must be tested just like any application.

Types of Testing:



Unit tests for tool calls.



Scenario tests for workflows.



Edge cases for unpredictable inputs.

Tools:



DeepEval



LangTest



Checkmate.

Tip

Treat LLM testing like QA—cover happy paths, failures, and adversarial cases.





Christopher R

DEPLOY & SCALE TO PRODUCTION

Prototyping is easy—scaling is hard.

Best Practices:

- ☐ Containerize agents for portability.
- ☐ Use version control for prompts + workflows.
- ☐ Set up monitoring for uptime + error handling.

Platforms:

-  Modal
-  BentoML
-  Vercel AI SDK
-  Fly.io

Tip

Start with a limited beta release to gather real-world usage data.





Christopher R

YOUR ROADMAP TO AI AGENTS

